

# aml-assignment-2

March 31, 2024

## Assignment 2: Convolution.

SHIVA CHIATANYA GOUD GADILA - 811252042

MALLENAHALLI VIJAYA PRANEETH SIMHA - 811252718

Retrieving the data

```
[1]: !mkdir ~/.kaggle  
      !cp kaggle.json ~/.kaggle/  
      !chmod 600 ~/.kaggle/kaggle.json
```

```
[2]: !kaggle competitions download -c dogs-vs-cats
```

```
Downloading dogs-vs-cats.zip to /content  
100% 811M/812M [00:38<00:00, 24.0MB/s]  
100% 812M/812M [00:39<00:00, 21.8MB/s]
```

```
[3]: !unzip -qq dogs-vs-cats.zip  
      !unzip -qq train.zip
```

Transferring the images to the training, validation, and test directories.

```
[4]: import os, shutil, pathlib  
  
original_dir = pathlib.Path("train")  
new_base_dir = pathlib.Path("cats_vs_dogs_small")  
  
def make_subset(subset_name, start_index, end_index):  
    for category in ("cat", "dog"):  
        dir = new_base_dir / subset_name / category  
        os.makedirs(dir)  
        fnames = [f"{category}.{i}.jpg" for i in range(start_index, end_index)]  
        for fname in fnames:  
            shutil.copyfile(src=original_dir / fname,  
                            dst=dir / fname)
```

TRAINING THE CONVNET NETWORK FROM SCRATCH:

MODEL 1: TRAINING SAMPLE OF 1000, VALIDATION SAMPLE OF 500 AND TEST SAMPLE OF 500

```
[5]: make_subset("test", start_index=0, end_index=500)
      make_subset("validation", start_index=500, end_index=1000)
      make_subset("train", start_index=1000, end_index=2000)
```

```
[6]: from tensorflow import keras
      from tensorflow.keras import layers
```

```
[7]: #Instantiating a small convnet for dogs vs. cats classification:

inputs = keras.Input(shape=(180, 180, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
Model_1 = keras.Model(inputs=inputs, outputs=outputs)
```

```
[8]: Model_1.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 180, 180, 3)]	0
rescaling (Rescaling)	(None, 180, 180, 3)	0
conv2d (Conv2D)	(None, 178, 178, 32)	896
max_pooling2d (MaxPooling2D)	(None, 89, 89, 32)	0
conv2d_1 (Conv2D)	(None, 87, 87, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 43, 43, 64)	0
conv2d_2 (Conv2D)	(None, 41, 41, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 20, 20, 128)	0

g2D)

conv2d_3 (Conv2D)	(None, 18, 18, 256)	295168
max_pooling2d_3 (MaxPoolin g2D)	(None, 9, 9, 256)	0
conv2d_4 (Conv2D)	(None, 7, 7, 256)	590080
flatten (Flatten)	(None, 12544)	0
dense (Dense)	(None, 1)	12545

```
=====
Total params: 991041 (3.78 MB)
Trainable params: 991041 (3.78 MB)
Non-trainable params: 0 (0.00 Byte)
-----
```

[9]: *#Configuring the model for training:*

```
Model_1.compile(loss="binary_crossentropy",
                 optimizer="rmsprop",
                 metrics=["accuracy"])
```

DATA PREPROCESSING:

[10]: *#Using image\_dataset\_from\_directory to read images*

```
from tensorflow.keras.utils import image_dataset_from_directory

train_dataset = image_dataset_from_directory(
    new_base_dir / "train",
    image_size=(180, 180),
    batch_size=32)
validation_dataset = image_dataset_from_directory(
    new_base_dir / "validation",
    image_size=(180, 180),
    batch_size=32)
test_dataset = image_dataset_from_directory(
    new_base_dir / "test",
    image_size=(180, 180),
    batch_size=32)
```

Found 2000 files belonging to 2 classes.  
Found 1000 files belonging to 2 classes.  
Found 1000 files belonging to 2 classes.

```
[11]: import numpy as np
import tensorflow as tf
random_numbers = np.random.normal(size=(1000, 16))
dataset = tf.data.Dataset.from_tensor_slices(random_numbers)
```

```
[12]: for i, element in enumerate(dataset):
    print(element.shape)
    if i >= 2:
        break
```

(16,)

(16,)

(16,)

```
[13]: batched_dataset = dataset.batch(32)
for i, element in enumerate(batched_dataset):
    print(element.shape)
    if i >= 2:
        break
```

(32, 16)

(32, 16)

(32, 16)

```
[14]: reshaped_dataset = dataset.map(lambda x: tf.reshape(x, (4, 4)))
for i, element in enumerate(reshaped_dataset):
    print(element.shape)
    if i >= 2:
        break
```

(4, 4)

(4, 4)

(4, 4)

```
[15]: #Displaying the shapes of the data and labels yielded by the Dataset:
```

```
for data_batch, labels_batch in train_dataset:
    print("data batch shape:", data_batch.shape)
    print("labels batch shape:", labels_batch.shape)
    break
```

data batch shape: (32, 180, 180, 3)

labels batch shape: (32,)

```
[16]: #Fitting the model using a Dataset
```

```
callbacks = [
    keras.callbacks.ModelCheckpoint(
```

```

        filepath="convnet_from_scratch.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = Model_1.fit(
    train_dataset,
    epochs=30,
    validation_data=validation_dataset,
    callbacks=callbacks)

```

Epoch 1/30

63/63 [=====] - 14s 133ms/step - loss: 0.6924 - accuracy: 0.5425 - val\_loss: 0.7394 - val\_accuracy: 0.5000

Epoch 2/30

63/63 [=====] - 5s 75ms/step - loss: 0.6884 - accuracy: 0.5610 - val\_loss: 0.6769 - val\_accuracy: 0.5710

Epoch 3/30

63/63 [=====] - 4s 63ms/step - loss: 0.6539 - accuracy: 0.6200 - val\_loss: 0.6455 - val\_accuracy: 0.6200

Epoch 4/30

63/63 [=====] - 6s 95ms/step - loss: 0.6224 - accuracy: 0.6550 - val\_loss: 0.6343 - val\_accuracy: 0.6480

Epoch 5/30

63/63 [=====] - 5s 74ms/step - loss: 0.5940 - accuracy: 0.6820 - val\_loss: 0.6235 - val\_accuracy: 0.6650

Epoch 6/30

63/63 [=====] - 4s 58ms/step - loss: 0.5722 - accuracy: 0.7195 - val\_loss: 0.7598 - val\_accuracy: 0.5520

Epoch 7/30

63/63 [=====] - 4s 64ms/step - loss: 0.5337 - accuracy: 0.7280 - val\_loss: 0.5439 - val\_accuracy: 0.7300

Epoch 8/30

63/63 [=====] - 7s 111ms/step - loss: 0.5034 - accuracy: 0.7595 - val\_loss: 0.5992 - val\_accuracy: 0.7000

Epoch 9/30

63/63 [=====] - 4s 58ms/step - loss: 0.4716 - accuracy: 0.7775 - val\_loss: 0.6081 - val\_accuracy: 0.7140

Epoch 10/30

63/63 [=====] - 5s 81ms/step - loss: 0.4234 - accuracy: 0.7980 - val\_loss: 0.6847 - val\_accuracy: 0.7080

Epoch 11/30

63/63 [=====] - 4s 61ms/step - loss: 0.3738 - accuracy: 0.8390 - val\_loss: 0.6834 - val\_accuracy: 0.6830

Epoch 12/30

63/63 [=====] - 5s 76ms/step - loss: 0.3177 - accuracy: 0.8590 - val\_loss: 0.6105 - val\_accuracy: 0.7450


Epoch 13/30


63/63 [=====] - 6s 87ms/step - loss: 0.2444 - accuracy: 0.8935 - val\_loss: 0.6725 - val\_accuracy: 0.7420  
Epoch 14/30  
63/63 [=====] - 4s 65ms/step - loss: 0.2137 - accuracy: 0.9210 - val\_loss: 0.8268 - val\_accuracy: 0.7450  
Epoch 15/30  
63/63 [=====] - 7s 111ms/step - loss: 0.1625 - accuracy: 0.9415 - val\_loss: 0.9790 - val\_accuracy: 0.7300  
Epoch 16/30  
63/63 [=====] - 4s 61ms/step - loss: 0.1498 - accuracy: 0.9500 - val\_loss: 0.8154 - val\_accuracy: 0.7390  
Epoch 17/30  
63/63 [=====] - 5s 82ms/step - loss: 0.0886 - accuracy: 0.9695 - val\_loss: 1.0547 - val\_accuracy: 0.7650  
Epoch 18/30  
63/63 [=====] - 6s 88ms/step - loss: 0.0947 - accuracy: 0.9685 - val\_loss: 1.0872 - val\_accuracy: 0.7100  
Epoch 19/30  
63/63 [=====] - 4s 59ms/step - loss: 0.0571 - accuracy: 0.9790 - val\_loss: 1.7071 - val\_accuracy: 0.6430  
Epoch 20/30  
63/63 [=====] - 4s 64ms/step - loss: 0.0766 - accuracy: 0.9745 - val\_loss: 1.3385 - val\_accuracy: 0.7330  
Epoch 21/30  
63/63 [=====] - 7s 97ms/step - loss: 0.0426 - accuracy: 0.9845 - val\_loss: 1.3004 - val\_accuracy: 0.7410  
Epoch 22/30  
63/63 [=====] - 4s 57ms/step - loss: 0.0671 - accuracy: 0.9775 - val\_loss: 1.4239 - val\_accuracy: 0.7530  
Epoch 23/30  
63/63 [=====] - 6s 95ms/step - loss: 0.0356 - accuracy: 0.9875 - val\_loss: 2.4992 - val\_accuracy: 0.7120  
Epoch 24/30  
63/63 [=====] - 5s 71ms/step - loss: 0.0698 - accuracy: 0.9820 - val\_loss: 1.4507 - val\_accuracy: 0.7310  
Epoch 25/30  
63/63 [=====] - 4s 57ms/step - loss: 0.0862 - accuracy: 0.9790 - val\_loss: 1.4374 - val\_accuracy: 0.7430  
Epoch 26/30  
63/63 [=====] - 5s 74ms/step - loss: 0.0278 - accuracy: 0.9890 - val\_loss: 1.5221 - val\_accuracy: 0.7550  
Epoch 27/30  
63/63 [=====] - 6s 94ms/step - loss: 0.0313 - accuracy: 0.9905 - val\_loss: 1.9706 - val\_accuracy: 0.7400  
Epoch 28/30  
63/63 [=====] - 4s 58ms/step - loss: 0.0444 - accuracy: 0.9860 - val\_loss: 1.9662 - val\_accuracy: 0.7370  
Epoch 29/30

```
63/63 [=====] - 4s 61ms/step - loss: 0.0191 - accuracy:
0.9945 - val_loss: 1.8982 - val_accuracy: 0.7420
Epoch 30/30
63/63 [=====] - 6s 98ms/step - loss: 0.0498 - accuracy:
0.9810 - val_loss: 1.6993 - val_accuracy: 0.7390
```

```
[17]: #Displaying curves of loss and accuracy during training:
```

```
import matplotlib.pyplot as plt
accuracy = history.history["accuracy"]
val_accuracy = history.history["val_accuracy"]
loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(accuracy) + 1)
plt.plot(epochs, accuracy, "bo", label="Training accuracy")
plt.plot(epochs, val_accuracy, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
plt.show()#
```





[18]: *#Evaluating the model on the test set:*

```
test_model = keras.models.load_model("convnet_from_scratch.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")
```

32/32 [=====] - 1s 29ms/step - loss: 0.5348 - accuracy: 0.7370

Test accuracy: 0.737

REDUCING THE OVERFITTING MODELS

MODEL 2- USING DATA AUGMENTATION:

```
[19]: from tensorflow import keras
      from tensorflow.keras import layers
      data_augmentation = keras.Sequential(
          [
              layers.RandomFlip("horizontal"),
              layers.RandomRotation(0.1),
```



```

        layers.RandomZoom(0.2),
    ]
)

```

```

[20]: inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = layers.Rescaling(1./255)(x)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
Model_2 = keras.Model(inputs=inputs, outputs=outputs)

Model_2.compile(loss="binary_crossentropy",
                 optimizer="rmsprop",
                 metrics=["accuracy"])

```

```

[21]: callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch_with_augmentation.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = Model_2.fit(
    train_dataset,
    epochs=30,
    validation_data=validation_dataset,
    callbacks=callbacks)

```

Epoch 1/30

63/63 [=====] - 10s 103ms/step - loss: 0.7007 - accuracy: 0.5030 - val\_loss: 0.6927 - val\_accuracy: 0.5000

Epoch 2/30

63/63 [=====] - 4s 64ms/step - loss: 0.6943 - accuracy: 0.5000 - val\_loss: 0.6909 - val\_accuracy: 0.5640

Epoch 3/30

63/63 [=====] - 4s 60ms/step - loss: 0.6893 - accuracy: 0.5365 - val\_loss: 0.6719 - val\_accuracy: 0.6260

Epoch 4/30

63/63 [=====] - 6s 98ms/step - loss: 0.6765 - accuracy:

0.5900 - val\_loss: 0.6565 - val\_accuracy: 0.6170  
Epoch 5/30  
63/63 [=====] - 4s 60ms/step - loss: 0.6600 - accuracy:  
0.6355 - val\_loss: 0.6616 - val\_accuracy: 0.6000  
Epoch 6/30  
63/63 [=====] - 6s 86ms/step - loss: 0.6358 - accuracy:  
0.6430 - val\_loss: 0.6384 - val\_accuracy: 0.6330  
Epoch 7/30  
63/63 [=====] - 6s 89ms/step - loss: 0.6434 - accuracy:  
0.6580 - val\_loss: 0.6226 - val\_accuracy: 0.6540  
Epoch 8/30  
63/63 [=====] - 4s 61ms/step - loss: 0.6025 - accuracy:  
0.6725 - val\_loss: 0.5940 - val\_accuracy: 0.6950  
Epoch 9/30  
63/63 [=====] - 4s 58ms/step - loss: 0.6046 - accuracy:  
0.6850 - val\_loss: 0.6440 - val\_accuracy: 0.6660  
Epoch 10/30  
63/63 [=====] - 7s 98ms/step - loss: 0.5990 - accuracy:  
0.6790 - val\_loss: 0.6212 - val\_accuracy: 0.7000  
Epoch 11/30  
63/63 [=====] - 4s 60ms/step - loss: 0.5847 - accuracy:  
0.6875 - val\_loss: 0.5835 - val\_accuracy: 0.6870  
Epoch 12/30  
63/63 [=====] - 7s 100ms/step - loss: 0.5803 -  
accuracy: 0.6965 - val\_loss: 0.5810 - val\_accuracy: 0.7010  
Epoch 13/30  
63/63 [=====] - 5s 69ms/step - loss: 0.5737 - accuracy:  
0.7075 - val\_loss: 0.5667 - val\_accuracy: 0.7170  
Epoch 14/30  
63/63 [=====] - 4s 61ms/step - loss: 0.5404 - accuracy:  
0.7310 - val\_loss: 0.5244 - val\_accuracy: 0.7410  
Epoch 15/30  
63/63 [=====] - 6s 86ms/step - loss: 0.5473 - accuracy:  
0.7315 - val\_loss: 0.5199 - val\_accuracy: 0.7560  
Epoch 16/30  
63/63 [=====] - 6s 87ms/step - loss: 0.5290 - accuracy:  
0.7295 - val\_loss: 0.5523 - val\_accuracy: 0.7380  
Epoch 17/30  
63/63 [=====] - 5s 76ms/step - loss: 0.5272 - accuracy:  
0.7415 - val\_loss: 0.5913 - val\_accuracy: 0.7260  
Epoch 18/30  
63/63 [=====] - 7s 101ms/step - loss: 0.5089 -  
accuracy: 0.7435 - val\_loss: 0.7099 - val\_accuracy: 0.6080  
Epoch 19/30  
63/63 [=====] - 4s 60ms/step - loss: 0.5016 - accuracy:  
0.7475 - val\_loss: 0.4918 - val\_accuracy: 0.7700  
Epoch 20/30  
63/63 [=====] - 4s 60ms/step - loss: 0.4940 - accuracy:

```

0.7620 - val_loss: 0.9062 - val_accuracy: 0.6680
Epoch 21/30
63/63 [=====] - 6s 87ms/step - loss: 0.4851 - accuracy:
0.7630 - val_loss: 0.4852 - val_accuracy: 0.7840
Epoch 22/30
63/63 [=====] - 6s 81ms/step - loss: 0.4713 - accuracy:
0.7765 - val_loss: 0.6031 - val_accuracy: 0.7510
Epoch 23/30
63/63 [=====] - 4s 58ms/step - loss: 0.4630 - accuracy:
0.7785 - val_loss: 0.5476 - val_accuracy: 0.7550
Epoch 24/30
63/63 [=====] - 4s 61ms/step - loss: 0.4627 - accuracy:
0.7815 - val_loss: 0.4785 - val_accuracy: 0.7820
Epoch 25/30
63/63 [=====] - 7s 102ms/step - loss: 0.4402 -
accuracy: 0.8010 - val_loss: 0.5763 - val_accuracy: 0.7380
Epoch 26/30
63/63 [=====] - 4s 63ms/step - loss: 0.4411 - accuracy:
0.8010 - val_loss: 0.5013 - val_accuracy: 0.7800
Epoch 27/30
63/63 [=====] - 4s 60ms/step - loss: 0.4409 - accuracy:
0.7935 - val_loss: 0.4463 - val_accuracy: 0.8140
Epoch 28/30
63/63 [=====] - 5s 73ms/step - loss: 0.4281 - accuracy:
0.7980 - val_loss: 0.4518 - val_accuracy: 0.8010
Epoch 29/30
63/63 [=====] - 6s 87ms/step - loss: 0.4113 - accuracy:
0.8160 - val_loss: 0.4863 - val_accuracy: 0.7820
Epoch 30/30
63/63 [=====] - 4s 58ms/step - loss: 0.4127 - accuracy:
0.8185 - val_loss: 0.5258 - val_accuracy: 0.7560

```

```

[22]: test_model = keras.models.load_model(
        "convnet_from_scratch_with_augmentation.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

```

```

32/32 [=====] - 2s 32ms/step - loss: 0.4246 - accuracy:
0.8190
Test accuracy: 0.819

```

```
[ ]:
```

```

[23]: #Defining a data augmentation stage to add to an image model:

```

```

data_augmentation = keras.Sequential(
    [

```

```
layers.RandomFlip("horizontal"),  
layers.RandomRotation(0.1),  
layers.RandomZoom(0.2),  
]  
)
```

[24]: *#Displaying some randomly augmented training images*

```
plt.figure(figsize=(10, 10))  
for images, _ in train_dataset.take(1):  
    for i in range(9):  
        augmented_images = data_augmentation(images)  
        ax = plt.subplot(3, 3, i + 1)  
        plt.imshow(augmented_images[0].numpy().astype("uint8"))  
        plt.axis("off")
```



### MODEL 3 - DROPOUT METHOD

```
[25]: inputs = keras.Input(shape=(180, 180, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
Model_3 = keras.Model(inputs=inputs, outputs=outputs)

Model_3.compile(loss="binary_crossentropy",
                 optimizer="rmsprop",
                 metrics=["accuracy"])
```

```
[26]: callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch_with_dropout.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = Model_3.fit(
    train_dataset,
    epochs=30,
    validation_data=validation_dataset,
    callbacks=callbacks)
```

Epoch 1/30

63/63 [=====] - 10s 102ms/step - loss: 0.7047 - accuracy: 0.5145 - val\_loss: 0.6921 - val\_accuracy: 0.5000

Epoch 2/30

63/63 [=====] - 4s 60ms/step - loss: 0.6967 - accuracy: 0.5205 - val\_loss: 0.6866 - val\_accuracy: 0.5730

Epoch 3/30

63/63 [=====] - 4s 59ms/step - loss: 0.6908 - accuracy: 0.5410 - val\_loss: 0.6922 - val\_accuracy: 0.5010

Epoch 4/30

63/63 [=====] - 5s 79ms/step - loss: 0.6767 - accuracy:

0.6070 - val\_loss: 0.6622 - val\_accuracy: 0.6150  
Epoch 5/30  
63/63 [=====] - 6s 89ms/step - loss: 0.6476 - accuracy:  
0.6345 - val\_loss: 0.6564 - val\_accuracy: 0.6210  
Epoch 6/30  
63/63 [=====] - 4s 59ms/step - loss: 0.6248 - accuracy:  
0.6610 - val\_loss: 0.6433 - val\_accuracy: 0.6300  
Epoch 7/30  
63/63 [=====] - 7s 110ms/step - loss: 0.5890 -  
accuracy: 0.6930 - val\_loss: 0.6144 - val\_accuracy: 0.6590  
Epoch 8/30  
63/63 [=====] - 4s 58ms/step - loss: 0.5901 - accuracy:  
0.6825 - val\_loss: 0.7780 - val\_accuracy: 0.6200  
Epoch 9/30  
63/63 [=====] - 4s 60ms/step - loss: 0.5602 - accuracy:  
0.7085 - val\_loss: 0.6061 - val\_accuracy: 0.6770  
Epoch 10/30  
63/63 [=====] - 6s 87ms/step - loss: 0.5277 - accuracy:  
0.7325 - val\_loss: 0.5595 - val\_accuracy: 0.7180  
Epoch 11/30  
63/63 [=====] - 6s 83ms/step - loss: 0.5016 - accuracy:  
0.7645 - val\_loss: 0.6106 - val\_accuracy: 0.6950  
Epoch 12/30  
63/63 [=====] - 4s 57ms/step - loss: 0.4562 - accuracy:  
0.7770 - val\_loss: 0.5892 - val\_accuracy: 0.6910  
Epoch 13/30  
63/63 [=====] - 4s 61ms/step - loss: 0.4259 - accuracy:  
0.7995 - val\_loss: 0.5687 - val\_accuracy: 0.7450  
Epoch 14/30  
63/63 [=====] - 7s 111ms/step - loss: 0.3854 -  
accuracy: 0.8265 - val\_loss: 0.5532 - val\_accuracy: 0.7380  
Epoch 15/30  
63/63 [=====] - 4s 61ms/step - loss: 0.3366 - accuracy:  
0.8520 - val\_loss: 0.6629 - val\_accuracy: 0.7450  
Epoch 16/30  
63/63 [=====] - 5s 68ms/step - loss: 0.2951 - accuracy:  
0.8720 - val\_loss: 0.7394 - val\_accuracy: 0.7280  
Epoch 17/30  
63/63 [=====] - 7s 99ms/step - loss: 0.2670 - accuracy:  
0.8845 - val\_loss: 0.6280 - val\_accuracy: 0.7300  
Epoch 18/30  
63/63 [=====] - 4s 59ms/step - loss: 0.2183 - accuracy:  
0.9125 - val\_loss: 0.7169 - val\_accuracy: 0.7580  
Epoch 19/30  
63/63 [=====] - 4s 61ms/step - loss: 0.1816 - accuracy:  
0.9245 - val\_loss: 1.0737 - val\_accuracy: 0.6700  
Epoch 20/30  
63/63 [=====] - 7s 101ms/step - loss: 0.1563 -

```

accuracy: 0.9420 - val_loss: 0.8744 - val_accuracy: 0.7420
Epoch 21/30
63/63 [=====] - 4s 62ms/step - loss: 0.1253 - accuracy:
0.9565 - val_loss: 1.1023 - val_accuracy: 0.7280
Epoch 22/30
63/63 [=====] - 4s 65ms/step - loss: 0.1093 - accuracy:
0.9605 - val_loss: 1.0383 - val_accuracy: 0.7550
Epoch 23/30
63/63 [=====] - 6s 90ms/step - loss: 0.1114 - accuracy:
0.9585 - val_loss: 0.9574 - val_accuracy: 0.7510
Epoch 24/30
63/63 [=====] - 4s 58ms/step - loss: 0.0992 - accuracy:
0.9680 - val_loss: 1.1496 - val_accuracy: 0.7740
Epoch 25/30
63/63 [=====] - 7s 103ms/step - loss: 0.0758 -
accuracy: 0.9710 - val_loss: 1.2123 - val_accuracy: 0.7410
Epoch 26/30
63/63 [=====] - 5s 68ms/step - loss: 0.0550 - accuracy:
0.9785 - val_loss: 1.4881 - val_accuracy: 0.7410
Epoch 27/30
63/63 [=====] - 4s 58ms/step - loss: 0.0526 - accuracy:
0.9805 - val_loss: 1.3756 - val_accuracy: 0.7450
Epoch 28/30
63/63 [=====] - 4s 68ms/step - loss: 0.0676 - accuracy:
0.9760 - val_loss: 1.3157 - val_accuracy: 0.7570
Epoch 29/30
63/63 [=====] - 6s 93ms/step - loss: 0.0743 - accuracy:
0.9740 - val_loss: 1.7772 - val_accuracy: 0.7400
Epoch 30/30
63/63 [=====] - 4s 59ms/step - loss: 0.0666 - accuracy:
0.9770 - val_loss: 1.4717 - val_accuracy: 0.7550

```

```

[27]: test_model = keras.models.load_model(
        "convnet_from_scratch_with_dropout.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

```

```

32/32 [=====] - 1s 28ms/step - loss: 0.5331 - accuracy:
0.7580
Test accuracy: 0.758

```

MODEL 4 - DATA AUGMENTATION AND DROPOUT METHOD:

```

[28]: data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
    ]
)

```

```
]
)
```

```
[29]: inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = layers.Rescaling(1./255)(x)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
Model_4 = keras.Model(inputs=inputs, outputs=outputs)

Model_4.compile(loss="binary_crossentropy",
                 optimizer="rmsprop",
                 metrics=["accuracy"])
```

```
[30]: callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch_with_augmentation_dropout.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = Model_4.fit(
    train_dataset,
    epochs=30,
    validation_data=validation_dataset,
    callbacks=callbacks)
```

Epoch 1/30

63/63 [=====] - 9s 88ms/step - loss: 0.6996 - accuracy: 0.4925 - val\_loss: 0.6928 - val\_accuracy: 0.5790

Epoch 2/30

63/63 [=====] - 4s 59ms/step - loss: 0.6991 - accuracy: 0.5060 - val\_loss: 0.6949 - val\_accuracy: 0.5000

Epoch 3/30

63/63 [=====] - 7s 105ms/step - loss: 0.6927 - accuracy: 0.5185 - val\_loss: 0.6902 - val\_accuracy: 0.5250

Epoch 4/30

63/63 [=====] - 4s 62ms/step - loss: 0.6862 - accuracy:



0.5890 - val\_loss: 0.6873 - val\_accuracy: 0.5150  
Epoch 5/30  
63/63 [=====] - 4s 64ms/step - loss: 0.6687 - accuracy:  
0.5905 - val\_loss: 0.6276 - val\_accuracy: 0.6630  
Epoch 6/30  
63/63 [=====] - 7s 105ms/step - loss: 0.6337 -  
accuracy: 0.6445 - val\_loss: 0.6243 - val\_accuracy: 0.6660  
Epoch 7/30  
63/63 [=====] - 4s 62ms/step - loss: 0.6330 - accuracy:  
0.6560 - val\_loss: 0.6132 - val\_accuracy: 0.6790  
Epoch 8/30  
63/63 [=====] - 6s 88ms/step - loss: 0.6110 - accuracy:  
0.6720 - val\_loss: 0.7306 - val\_accuracy: 0.6030  
Epoch 9/30  
63/63 [=====] - 6s 87ms/step - loss: 0.5976 - accuracy:  
0.6820 - val\_loss: 0.5970 - val\_accuracy: 0.6800  
Epoch 10/30  
63/63 [=====] - 4s 59ms/step - loss: 0.5941 - accuracy:  
0.6825 - val\_loss: 0.6032 - val\_accuracy: 0.6760  
Epoch 11/30  
63/63 [=====] - 4s 65ms/step - loss: 0.5952 - accuracy:  
0.6895 - val\_loss: 0.6653 - val\_accuracy: 0.6330  
Epoch 12/30  
63/63 [=====] - 8s 114ms/step - loss: 0.5765 -  
accuracy: 0.7080 - val\_loss: 0.6308 - val\_accuracy: 0.6560  
Epoch 13/30  
63/63 [=====] - 4s 62ms/step - loss: 0.5711 - accuracy:  
0.7115 - val\_loss: 0.6257 - val\_accuracy: 0.6680  
Epoch 14/30  
63/63 [=====] - 6s 98ms/step - loss: 0.5556 - accuracy:  
0.7300 - val\_loss: 0.5580 - val\_accuracy: 0.7120  
Epoch 15/30  
63/63 [=====] - 5s 71ms/step - loss: 0.5391 - accuracy:  
0.7280 - val\_loss: 0.5752 - val\_accuracy: 0.7050  
Epoch 16/30  
63/63 [=====] - 5s 74ms/step - loss: 0.5421 - accuracy:  
0.7310 - val\_loss: 0.5564 - val\_accuracy: 0.7320  
Epoch 17/30  
63/63 [=====] - 7s 100ms/step - loss: 0.5264 -  
accuracy: 0.7375 - val\_loss: 0.5501 - val\_accuracy: 0.7270  
Epoch 18/30  
63/63 [=====] - 4s 61ms/step - loss: 0.5257 - accuracy:  
0.7425 - val\_loss: 0.5098 - val\_accuracy: 0.7470  
Epoch 19/30  
63/63 [=====] - 4s 61ms/step - loss: 0.5178 - accuracy:  
0.7450 - val\_loss: 0.5056 - val\_accuracy: 0.7590  
Epoch 20/30  
63/63 [=====] - 6s 95ms/step - loss: 0.5108 - accuracy:

```

0.7535 - val_loss: 0.5252 - val_accuracy: 0.7390
Epoch 21/30
63/63 [=====] - 4s 59ms/step - loss: 0.5104 - accuracy:
0.7530 - val_loss: 0.5707 - val_accuracy: 0.6960
Epoch 22/30
63/63 [=====] - 5s 84ms/step - loss: 0.5002 - accuracy:
0.7585 - val_loss: 0.6116 - val_accuracy: 0.6720
Epoch 23/30
63/63 [=====] - 6s 90ms/step - loss: 0.4796 - accuracy:
0.7755 - val_loss: 0.5508 - val_accuracy: 0.7300
Epoch 24/30
63/63 [=====] - 4s 61ms/step - loss: 0.5021 - accuracy:
0.7625 - val_loss: 0.4857 - val_accuracy: 0.7620
Epoch 25/30
63/63 [=====] - 4s 62ms/step - loss: 0.4656 - accuracy:
0.7850 - val_loss: 0.5184 - val_accuracy: 0.7560
Epoch 26/30
63/63 [=====] - 7s 102ms/step - loss: 0.4456 -
accuracy: 0.7945 - val_loss: 0.4910 - val_accuracy: 0.7600
Epoch 27/30
63/63 [=====] - 4s 60ms/step - loss: 0.4371 - accuracy:
0.8055 - val_loss: 0.5462 - val_accuracy: 0.7620
Epoch 28/30
63/63 [=====] - 4s 61ms/step - loss: 0.4358 - accuracy:
0.8010 - val_loss: 0.4393 - val_accuracy: 0.8010
Epoch 29/30
63/63 [=====] - 7s 103ms/step - loss: 0.4318 -
accuracy: 0.7990 - val_loss: 0.4653 - val_accuracy: 0.7780
Epoch 30/30
63/63 [=====] - 5s 68ms/step - loss: 0.4278 - accuracy:
0.8030 - val_loss: 0.4994 - val_accuracy: 0.7780

```

```

[31]: test_model = keras.models.load_model(
        "convnet_from_scratch_with_augmentation_dropout.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

```

```

32/32 [=====] - 1s 32ms/step - loss: 0.4004 - accuracy:
0.8180
Test accuracy: 0.818

```

MODEL 5 - INCREASING THE TRAINING SAMPLE SIZE TO 5000, INCLUDING MAXPOOLING, DATA AUGMENTATION AND DROPOUT TECHNIQUE (DROPOUT RATE = 0.05)

```

[32]: from tensorflow.keras.utils import image_dataset_from_directory

make_subset("train_1", start_index=0, end_index=5000)
make_subset("validation_1", start_index=5000, end_index=5500)

```

```

make_subset("test_1", start_index=5500, end_index=6000)

train_dataset_1 = image_dataset_from_directory(
    new_base_dir / "train_1",
    image_size=(180, 180),
    batch_size=32)
validation_dataset_1 = image_dataset_from_directory(
    new_base_dir / "validation_1",
    image_size=(180, 180),
    batch_size=32)
test_dataset_1 = image_dataset_from_directory(
    new_base_dir / "test_1",
    image_size=(180, 180),
    batch_size=32)

```

Found 10000 files belonging to 2 classes.  
 Found 1000 files belonging to 2 classes.  
 Found 1000 files belonging to 2 classes.

[33]: *#Defining a new convnet that includes image augmentation and dropout*

```

inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
Model_5 = keras.Model(inputs=inputs, outputs=outputs)

Model_5.compile(loss="binary_crossentropy",
                optimizer="rmsprop",
                metrics=["accuracy"])

```

[34]:

```

from tensorflow import keras
from tensorflow.keras import layers
import matplotlib.pyplot as plt
from keras.callbacks import EarlyStopping
from keras import regularizers

```

```
early_stopping_monitor = EarlyStopping(patience=10)
```

```
[35]: data_augmentation = keras.Sequential(  
    [  
        layers.RandomFlip("horizontal"),  
        layers.RandomRotation(0.1),  
        layers.RandomZoom(0.2),  
    ]  
)
```

```
[36]: plt.figure(figsize=(10, 10))  
for images, _ in train_dataset.take(1):  
    for i in range(9):  
        augmented_images = data_augmentation(images)  
        ax = plt.subplot(3, 3, i + 1)  
        plt.imshow(augmented_images[0].numpy().astype("uint8"))  
        plt.axis("off")
```



```
[37]: callbacks = [  
    keras.callbacks.ModelCheckpoint(  
        filepath="convnet_from_scratch.keras",  
        save_best_only=True,  
        monitor="val_loss"), early_stopping_monitor  
]  
history = Model_5.fit(  
    train_dataset_1,  
    epochs=30,  
    validation_data=validation_dataset,  
    callbacks=callbacks)
```

Epoch 1/30

313/313 [=====] - 24s 69ms/step - loss: 0.6958 -  
accuracy: 0.5484 - val\_loss: 0.6781 - val\_accuracy: 0.5300  
Epoch 2/30  
313/313 [=====] - 20s 61ms/step - loss: 0.6191 -  
accuracy: 0.6621 - val\_loss: 0.5799 - val\_accuracy: 0.7030  
Epoch 3/30  
313/313 [=====] - 17s 54ms/step - loss: 0.5420 -  
accuracy: 0.7266 - val\_loss: 0.5537 - val\_accuracy: 0.6980  
Epoch 4/30  
313/313 [=====] - 18s 56ms/step - loss: 0.4751 -  
accuracy: 0.7779 - val\_loss: 0.4254 - val\_accuracy: 0.8050  
Epoch 5/30  
313/313 [=====] - 18s 56ms/step - loss: 0.4124 -  
accuracy: 0.8146 - val\_loss: 0.3236 - val\_accuracy: 0.8640  
Epoch 6/30  
313/313 [=====] - 18s 57ms/step - loss: 0.3629 -  
accuracy: 0.8415 - val\_loss: 0.3336 - val\_accuracy: 0.8760  
Epoch 7/30  
313/313 [=====] - 18s 56ms/step - loss: 0.3144 -  
accuracy: 0.8658 - val\_loss: 0.2220 - val\_accuracy: 0.9060  
Epoch 8/30  
313/313 [=====] - 18s 57ms/step - loss: 0.2650 -  
accuracy: 0.8862 - val\_loss: 0.2131 - val\_accuracy: 0.9140  
Epoch 9/30  
313/313 [=====] - 18s 56ms/step - loss: 0.2324 -  
accuracy: 0.9041 - val\_loss: 0.1184 - val\_accuracy: 0.9550  
Epoch 10/30  
313/313 [=====] - 18s 55ms/step - loss: 0.1816 -  
accuracy: 0.9260 - val\_loss: 0.1209 - val\_accuracy: 0.9580  
Epoch 11/30  
313/313 [=====] - 18s 57ms/step - loss: 0.1597 -  
accuracy: 0.9377 - val\_loss: 0.1316 - val\_accuracy: 0.9500  
Epoch 12/30  
313/313 [=====] - 19s 59ms/step - loss: 0.1248 -  
accuracy: 0.9522 - val\_loss: 0.0904 - val\_accuracy: 0.9630  
Epoch 13/30  
313/313 [=====] - 18s 57ms/step - loss: 0.1196 -  
accuracy: 0.9560 - val\_loss: 0.0893 - val\_accuracy: 0.9640  
Epoch 14/30  
313/313 [=====] - 18s 58ms/step - loss: 0.1066 -  
accuracy: 0.9618 - val\_loss: 0.1361 - val\_accuracy: 0.9570  
Epoch 15/30  
313/313 [=====] - 18s 57ms/step - loss: 0.1022 -  
accuracy: 0.9640 - val\_loss: 0.1278 - val\_accuracy: 0.9540  
Epoch 16/30  
313/313 [=====] - 18s 57ms/step - loss: 0.0942 -  
accuracy: 0.9685 - val\_loss: 0.0360 - val\_accuracy: 0.9880  
Epoch 17/30

```

313/313 [=====] - 18s 57ms/step - loss: 0.0903 -
accuracy: 0.9687 - val_loss: 0.0272 - val_accuracy: 0.9900
Epoch 18/30
313/313 [=====] - 18s 57ms/step - loss: 0.0834 -
accuracy: 0.9712 - val_loss: 0.0432 - val_accuracy: 0.9870
Epoch 19/30
313/313 [=====] - 18s 57ms/step - loss: 0.0787 -
accuracy: 0.9743 - val_loss: 0.0241 - val_accuracy: 0.9910
Epoch 20/30
313/313 [=====] - 18s 57ms/step - loss: 0.0792 -
accuracy: 0.9747 - val_loss: 0.0185 - val_accuracy: 0.9940
Epoch 21/30
313/313 [=====] - 18s 57ms/step - loss: 0.0882 -
accuracy: 0.9730 - val_loss: 0.0174 - val_accuracy: 0.9940
Epoch 22/30
313/313 [=====] - 18s 56ms/step - loss: 0.0717 -
accuracy: 0.9781 - val_loss: 0.0186 - val_accuracy: 0.9940
Epoch 23/30
313/313 [=====] - 18s 56ms/step - loss: 0.0753 -
accuracy: 0.9767 - val_loss: 0.1112 - val_accuracy: 0.9720
Epoch 24/30
313/313 [=====] - 18s 57ms/step - loss: 0.0829 -
accuracy: 0.9769 - val_loss: 0.0190 - val_accuracy: 0.9910
Epoch 25/30
313/313 [=====] - 18s 57ms/step - loss: 0.0810 -
accuracy: 0.9783 - val_loss: 0.0216 - val_accuracy: 0.9920
Epoch 26/30
313/313 [=====] - 19s 58ms/step - loss: 0.0862 -
accuracy: 0.9761 - val_loss: 0.1217 - val_accuracy: 0.9790
Epoch 27/30
313/313 [=====] - 17s 54ms/step - loss: 0.0878 -
accuracy: 0.9776 - val_loss: 0.0245 - val_accuracy: 0.9900
Epoch 28/30
313/313 [=====] - 18s 57ms/step - loss: 0.0847 -
accuracy: 0.9791 - val_loss: 0.3962 - val_accuracy: 0.9350
Epoch 29/30
313/313 [=====] - 18s 57ms/step - loss: 0.0756 -
accuracy: 0.9794 - val_loss: 0.0177 - val_accuracy: 0.9930
Epoch 30/30
313/313 [=====] - 18s 56ms/step - loss: 0.0826 -
accuracy: 0.9796 - val_loss: 0.0750 - val_accuracy: 0.9830

```

```


[38]: accuracy = history.history["accuracy"]
      val_accuracy = history.history["val_accuracy"]
      loss = history.history["loss"]
      val_loss = history.history["val_loss"]
      epochs = range(1, len(accuracy) + 1)

```


```

plt.plot(epochs, accuracy, "bo", label="Training accuracy")
plt.plot(epochs, val_accuracy, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
plt.show()

```







```
[39]: test_model = keras.models.load_model("convnet_from_scratch.keras")
test_loss, test_acc = test_model.evaluate(test_dataset_1)
print(f"Test accuracy: {test_acc:.3f}")
```

```
32/32 [=====] - 2s 30ms/step - loss: 0.6666 - accuracy: 0.8870
```

```
Test accuracy: 0.887
```

MODEL 6 - INCREASING THE TRAINING SAMPLE TO 10000

```
[40]: from tensorflow.keras.utils import image_dataset_from_directory

make_subset("train_4", start_index=0, end_index=10000)
make_subset("validation_4", start_index=10000, end_index=10500)
make_subset("test_4", start_index=10500, end_index=11000)

train_dataset_4 = image_dataset_from_directory(
    new_base_dir / "train_4",
    image_size=(180, 180),
    batch_size=32)
validation_dataset_4 = image_dataset_from_directory(
```

```

new_base_dir / "validation_4",
image_size=(180, 180),
batch_size=32)
test_dataset_4 = image_dataset_from_directory(
new_base_dir / "test_4",
image_size=(180, 180),
batch_size=32)

```

Found 20000 files belonging to 2 classes.  
Found 1000 files belonging to 2 classes.  
Found 1000 files belonging to 2 classes.

```

[41]: inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
Model_6 = keras.Model(inputs=inputs, outputs=outputs)

Model_6.compile(loss="binary_crossentropy",
optimizer="rmsprop",
metrics=["accuracy"])

```

```

[42]: data_augmentation = keras.Sequential(
[
layers.RandomFlip("horizontal"),
layers.RandomRotation(0.1),
layers.RandomZoom(0.2),
]
)

```

```

[43]: plt.figure(figsize=(10, 10))
for images, _ in train_dataset.take(1):
for i in range(9):
augmented_images = data_augmentation(images)
ax = plt.subplot(3, 3, i + 1)
plt.imshow(augmented_images[0].numpy().astype("uint8"))

```

```
plt.axis("off")
```



```
[44]: callbacks = [  
    keras.callbacks.ModelCheckpoint(  
        filepath="convnet_from_scratch.keras",  
        save_best_only=True,  
        monitor="val_loss"), early_stopping_monitor  
]  
history = Model_6.fit(  
    train_dataset_4,  
    epochs=30,  
    validation_data=validation_dataset_4,  
    callbacks=callbacks)
```

Epoch 1/30  
625/625 [=====] - 36s 52ms/step - loss: 0.6409 -  
accuracy: 0.6166 - val\_loss: 0.5234 - val\_accuracy: 0.7420  
Epoch 2/30  
625/625 [=====] - 34s 54ms/step - loss: 0.4952 -  
accuracy: 0.7580 - val\_loss: 0.4023 - val\_accuracy: 0.8050  
Epoch 3/30  
625/625 [=====] - 34s 54ms/step - loss: 0.3844 -  
accuracy: 0.8268 - val\_loss: 0.5215 - val\_accuracy: 0.7410  
Epoch 4/30  
625/625 [=====] - 34s 55ms/step - loss: 0.3033 -  
accuracy: 0.8701 - val\_loss: 0.4296 - val\_accuracy: 0.7970  
Epoch 5/30  
625/625 [=====] - 35s 55ms/step - loss: 0.2507 -  
accuracy: 0.8957 - val\_loss: 0.2320 - val\_accuracy: 0.8970  
Epoch 6/30  
625/625 [=====] - 34s 55ms/step - loss: 0.2168 -  
accuracy: 0.9107 - val\_loss: 0.3164 - val\_accuracy: 0.8970  
Epoch 7/30  
625/625 [=====] - 35s 55ms/step - loss: 0.1805 -  
accuracy: 0.9277 - val\_loss: 0.2794 - val\_accuracy: 0.9120  
Epoch 8/30  
625/625 [=====] - 34s 55ms/step - loss: 0.1611 -  
accuracy: 0.9373 - val\_loss: 0.2463 - val\_accuracy: 0.9040  
Epoch 9/30  
625/625 [=====] - 34s 55ms/step - loss: 0.1506 -  
accuracy: 0.9417 - val\_loss: 0.2284 - val\_accuracy: 0.9230  
Epoch 10/30  
625/625 [=====] - 35s 55ms/step - loss: 0.1425 -  
accuracy: 0.9473 - val\_loss: 0.2410 - val\_accuracy: 0.9050  
Epoch 11/30  
625/625 [=====] - 33s 53ms/step - loss: 0.1309 -  
accuracy: 0.9523 - val\_loss: 0.2569 - val\_accuracy: 0.9170  
Epoch 12/30  
625/625 [=====] - 34s 53ms/step - loss: 0.1258 -  
accuracy: 0.9549 - val\_loss: 0.2676 - val\_accuracy: 0.9260  
Epoch 13/30  
625/625 [=====] - 34s 55ms/step - loss: 0.1228 -  
accuracy: 0.9595 - val\_loss: 0.3248 - val\_accuracy: 0.9160  
Epoch 14/30  
625/625 [=====] - 34s 53ms/step - loss: 0.1209 -  
accuracy: 0.9602 - val\_loss: 0.4604 - val\_accuracy: 0.8910  
Epoch 15/30  
625/625 [=====] - 34s 55ms/step - loss: 0.1124 -  
accuracy: 0.9642 - val\_loss: 0.6580 - val\_accuracy: 0.8740  
Epoch 16/30  
625/625 [=====] - 37s 58ms/step - loss: 0.1231 -  
accuracy: 0.9612 - val\_loss: 0.3833 - val\_accuracy: 0.9160

Epoch 17/30

625/625 [=====] - 36s 58ms/step - loss: 0.1246 - accuracy: 0.9650 - val\_loss: 0.5290 - val\_accuracy: 0.8860


Epoch 18/30


625/625 [=====] - 37s 59ms/step - loss: 0.1229 - accuracy: 0.9649 - val\_loss: 0.3659 - val\_accuracy: 0.9220

Epoch 19/30

625/625 [=====] - 35s 56ms/step - loss: 0.1294 - accuracy: 0.9654 - val\_loss: 0.3923 - val\_accuracy: 0.8990

```
[45]: accuracy = history.history["accuracy"]
val_accuracy = history.history["val_accuracy"]
loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(accracy) + 1)
plt.plot(epochs, accracy, "bo", label="Training accracy")
plt.plot(epochs, val_accuracy, "b", label="Validation accracy")
plt.title("Training and validation accracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
plt.show()
```





```
[46]: test_model = keras.models.load_model("convnet_from_scratch.keras")
test_loss, test_acc = test_model.evaluate(test_dataset_4)
print(f"Test accuracy: {test_acc:.3f}")
```

```
32/32 [=====] - 1s 32ms/step - loss: 0.3237 - accuracy: 0.8830
```

```
Test accuracy: 0.883
```

INSTANTIATING AND FREEZING THE VGG16 CONVOLUTIONAL BASE

PRE-TRAINED MODEL

MODEL 7 - SAMPLE SIZE OF 1000

```
[79]: import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers

# Load pre-trained VGG16 model without the top layer
conv_base = keras.applications.VGG16(
    weights="imagenet",
    include_top=False,
    input_shape=(180, 180, 3))
```

```

# Freeze convolutional base layers
conv_base.trainable = True
for layer in conv_base.layers[:-4]:
    layer.trainable = False

# Define data augmentation pipeline
data_augmentation = keras.Sequential([
    layers.RandomFlip("horizontal"),
    layers.RandomRotation(0.1),
    layers.RandomZoom(0.2),
])

# Define your model
inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = keras.applications.vgg16.preprocess_input(x)
x = conv_base(x)
x = layers.Flatten()(x)
x = layers.Dense(256, activation='relu')(x) # Add activation function
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)

# Compile the model
model.compile(loss="binary_crossentropy",
              optimizer=keras.optimizers.RMSprop(learning_rate=1e-5),
              metrics=["accuracy"])

# Define callbacks
callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="fine_tuning.h5",
        save_best_only=True,
        monitor="val_loss")
]

# Train the model
history = model.fit(
    train_dataset,
    epochs=30,
    validation_data=validation_dataset,
    callbacks=callbacks
)

```

Epoch 1/30

63/63 [=====] - ETA: 0s - loss: 3.1088 - accuracy:



0.6780

/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103:  
UserWarning: You are saving your model as an HDF5 file via `model.save()`. This  
file format is considered legacy. We recommend using instead the native Keras  
format, e.g. `model.save('my\_model.keras')`.

saving\_api.save\_model(

63/63 [=====] - 14s 187ms/step - loss: 3.1088 -  
accuracy: 0.6780 - val\_loss: 0.5563 - val\_accuracy: 0.8880

Epoch 2/30

63/63 [=====] - 14s 216ms/step - loss: 0.9146 -  
accuracy: 0.8310 - val\_loss: 0.3115 - val\_accuracy: 0.9330

Epoch 3/30

63/63 [=====] - 12s 180ms/step - loss: 0.5377 -  
accuracy: 0.8710 - val\_loss: 0.2285 - val\_accuracy: 0.9380

Epoch 4/30

63/63 [=====] - 11s 179ms/step - loss: 0.3763 -  
accuracy: 0.9050 - val\_loss: 0.1930 - val\_accuracy: 0.9510

Epoch 5/30

63/63 [=====] - 11s 169ms/step - loss: 0.2471 -  
accuracy: 0.9260 - val\_loss: 0.1460 - val\_accuracy: 0.9540

Epoch 6/30

63/63 [=====] - 11s 173ms/step - loss: 0.2122 -  
accuracy: 0.9320 - val\_loss: 0.1474 - val\_accuracy: 0.9600

Epoch 7/30

63/63 [=====] - 13s 198ms/step - loss: 0.1988 -  
accuracy: 0.9390 - val\_loss: 0.1725 - val\_accuracy: 0.9610

Epoch 8/30

63/63 [=====] - 12s 176ms/step - loss: 0.1541 -  
accuracy: 0.9495 - val\_loss: 0.1392 - val\_accuracy: 0.9660

Epoch 9/30

63/63 [=====] - 11s 166ms/step - loss: 0.1194 -  
accuracy: 0.9560 - val\_loss: 0.1691 - val\_accuracy: 0.9700

Epoch 10/30

63/63 [=====] - 11s 167ms/step - loss: 0.0983 -  
accuracy: 0.9685 - val\_loss: 0.1752 - val\_accuracy: 0.9690

Epoch 11/30

63/63 [=====] - 11s 175ms/step - loss: 0.0949 -  
accuracy: 0.9565 - val\_loss: 0.1763 - val\_accuracy: 0.9720

Epoch 12/30

63/63 [=====] - 11s 165ms/step - loss: 0.0917 -  
accuracy: 0.9680 - val\_loss: 0.1689 - val\_accuracy: 0.9730

Epoch 13/30

63/63 [=====] - 11s 170ms/step - loss: 0.0908 -  
accuracy: 0.9660 - val\_loss: 0.1488 - val\_accuracy: 0.9750

Epoch 14/30

63/63 [=====] - 11s 166ms/step - loss: 0.0644 -  
accuracy: 0.9760 - val\_loss: 0.1711 - val\_accuracy: 0.9750

Epoch 15/30  
63/63 [=====] - 11s 172ms/step - loss: 0.0634 -  
accuracy: 0.9750 - val\_loss: 0.1955 - val\_accuracy: 0.9720  
Epoch 16/30  
63/63 [=====] - 11s 172ms/step - loss: 0.0730 -  
accuracy: 0.9750 - val\_loss: 0.1890 - val\_accuracy: 0.9760  
Epoch 17/30  
63/63 [=====] - 11s 168ms/step - loss: 0.0620 -  
accuracy: 0.9785 - val\_loss: 0.2087 - val\_accuracy: 0.9740  
Epoch 18/30  
63/63 [=====] - 11s 167ms/step - loss: 0.0516 -  
accuracy: 0.9775 - val\_loss: 0.1962 - val\_accuracy: 0.9750  
Epoch 19/30  
63/63 [=====] - 11s 174ms/step - loss: 0.0712 -  
accuracy: 0.9800 - val\_loss: 0.1875 - val\_accuracy: 0.9780  
Epoch 20/30  
63/63 [=====] - 13s 196ms/step - loss: 0.0341 -  
accuracy: 0.9865 - val\_loss: 0.2172 - val\_accuracy: 0.9770  
Epoch 21/30  
63/63 [=====] - 11s 169ms/step - loss: 0.0416 -  
accuracy: 0.9860 - val\_loss: 0.2199 - val\_accuracy: 0.9750  
Epoch 22/30  
63/63 [=====] - 11s 171ms/step - loss: 0.0322 -  
accuracy: 0.9890 - val\_loss: 0.2049 - val\_accuracy: 0.9750  
Epoch 23/30  
63/63 [=====] - 11s 167ms/step - loss: 0.0355 -  
accuracy: 0.9860 - val\_loss: 0.2074 - val\_accuracy: 0.9780  
Epoch 24/30  
63/63 [=====] - 11s 165ms/step - loss: 0.0279 -  
accuracy: 0.9905 - val\_loss: 0.1995 - val\_accuracy: 0.9770  
Epoch 25/30  
63/63 [=====] - 11s 173ms/step - loss: 0.0182 -  
accuracy: 0.9930 - val\_loss: 0.2368 - val\_accuracy: 0.9780  
Epoch 26/30  
63/63 [=====] - 12s 194ms/step - loss: 0.0263 -  
accuracy: 0.9910 - val\_loss: 0.2054 - val\_accuracy: 0.9800  
Epoch 27/30  
63/63 [=====] - 11s 171ms/step - loss: 0.0320 -  
accuracy: 0.9900 - val\_loss: 0.1906 - val\_accuracy: 0.9780  
Epoch 28/30  
63/63 [=====] - 13s 197ms/step - loss: 0.0170 -  
accuracy: 0.9950 - val\_loss: 0.2366 - val\_accuracy: 0.9770  
Epoch 29/30  
63/63 [=====] - 11s 170ms/step - loss: 0.0337 -  
accuracy: 0.9880 - val\_loss: 0.2344 - val\_accuracy: 0.9750  
Epoch 30/30  
63/63 [=====] - 11s 167ms/step - loss: 0.0243 -  
accuracy: 0.9930 - val\_loss: 0.2212 - val\_accuracy: 0.9800

```
[80]: # Load the saved model
test_model = keras.models.load_model("fine_tuning.h5")

# Evaluate the model on the test dataset
test_loss, test_acc = test_model.evaluate(test_dataset)

# Print the test accuracy
print(f"Test accuracy: {test_acc:.3f}")
```

32/32 [=====] - 4s 94ms/step - loss: 0.1539 - accuracy: 0.9710

Test accuracy: 0.971

MODEL 8 - SAMPLE SIZE OF 5000

```
[81]: conv_base = keras.applications.vgg16.VGG16(
        weights="imagenet",
        include_top=False,
        input_shape=(180, 180, 3))
```

```
[82]: conv_base = keras.applications.vgg16.VGG16(
        weights="imagenet",
        include_top=False)

conv_base.trainable = True
for layer in conv_base.layers[:-4]:
    layer.trainable = False
```

```
[83]: data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
    ]
)
```

```
[84]: inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = keras.applications.vgg16.preprocess_input(x)
x = conv_base(x)
x = layers.Flatten()(x)
x = layers.Dense(256)(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
Model_8 = keras.Model(inputs, outputs)
Model_8.compile(loss="binary_crossentropy",
                optimizer=keras.optimizers.RMSprop(learning_rate=1e-5),
                metrics=["accuracy"])
```

```
[86]: # Define the callbacks
callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="fine_tuning2.h5",
        save_best_only=True,
        monitor="val_loss")
]

# Train Model_8
history = Model_8.fit(
    train_dataset_1,
    epochs=10,
    validation_data=validation_dataset_1,
    callbacks=callbacks
)
```

```
Epoch 1/10
313/313 [=====] - 42s 132ms/step - loss: 0.1885 -
accuracy: 0.9370 - val_loss: 0.1056 - val_accuracy: 0.9700
Epoch 2/10
313/313 [=====] - 39s 125ms/step - loss: 0.1466 -
accuracy: 0.9528 - val_loss: 0.1224 - val_accuracy: 0.9700
Epoch 3/10
313/313 [=====] - 40s 127ms/step - loss: 0.1133 -
accuracy: 0.9598 - val_loss: 0.1275 - val_accuracy: 0.9760
Epoch 4/10
313/313 [=====] - 41s 128ms/step - loss: 0.0931 -
accuracy: 0.9695 - val_loss: 0.0996 - val_accuracy: 0.9760
Epoch 5/10
313/313 [=====] - 40s 127ms/step - loss: 0.0935 -
accuracy: 0.9679 - val_loss: 0.1162 - val_accuracy: 0.9770
Epoch 6/10
313/313 [=====] - 42s 132ms/step - loss: 0.0842 -
accuracy: 0.9699 - val_loss: 0.1113 - val_accuracy: 0.9730
Epoch 7/10
313/313 [=====] - 42s 131ms/step - loss: 0.0725 -
accuracy: 0.9754 - val_loss: 0.0903 - val_accuracy: 0.9790
Epoch 8/10
313/313 [=====] - 41s 129ms/step - loss: 0.0651 -
accuracy: 0.9759 - val_loss: 0.1203 - val_accuracy: 0.9770
Epoch 9/10
313/313 [=====] - 39s 125ms/step - loss: 0.0578 -
accuracy: 0.9811 - val_loss: 0.1987 - val_accuracy: 0.9780
Epoch 10/10
313/313 [=====] - 40s 126ms/step - loss: 0.0652 -
accuracy: 0.9801 - val_loss: 0.1305 - val_accuracy: 0.9770
```

```
[87]: # Load the saved model
test_model = keras.models.load_model("fine_tuning2.h5")

# Evaluate the model on the test dataset
test_loss, test_acc = test_model.evaluate(test_dataset_1)

# Print the test accuracy
print(f"Test accuracy: {test_acc:.3f}")
```

32/32 [=====] - 4s 96ms/step - loss: 0.0723 - accuracy: 0.9810

Test accuracy: 0.981

MODEL 9 - SAMPLE SIZE OF 10000

```
[89]: conv_base = keras.applications.vgg16.VGG16(
        weights="imagenet",
        include_top=False,
        input_shape=(180, 180, 3))
```

```
[90]: conv_base = keras.applications.vgg16.VGG16(
        weights="imagenet",
        include_top=False)

conv_base.trainable = True
for layer in conv_base.layers[:-4]:
    layer.trainable = False
```

```
[91]: data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
    ]
)
```


```
[92]: inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = keras.applications.vgg16.preprocess_input(x)
x = conv_base(x)
x = layers.Flatten()(x)
x = layers.Dense(256)(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
Model_9 = keras.Model(inputs, outputs)
Model_9.compile(loss="binary_crossentropy",
                optimizer=keras.optimizers.RMSprop(learning_rate=1e-5),
                metrics=["accuracy"])
```


```
[94]: # Define the callbacks
callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="fine_tuning3.h5",
        save_best_only=True,
        monitor="val_loss")
]

# Train Model_9
history = Model_9.fit(
    train_dataset_4,
    epochs=10,
    validation_data=validation_dataset_4,
    callbacks=callbacks
)
```

```
Epoch 1/10
625/625 [=====] - 77s 123ms/step - loss: 0.1364 -
accuracy: 0.9528 - val_loss: 0.1705 - val_accuracy: 0.9600
Epoch 2/10
625/625 [=====] - 77s 123ms/step - loss: 0.1156 -
accuracy: 0.9626 - val_loss: 0.1063 - val_accuracy: 0.9700
Epoch 3/10
625/625 [=====] - 76s 121ms/step - loss: 0.0998 -
accuracy: 0.9679 - val_loss: 0.1021 - val_accuracy: 0.9720
Epoch 4/10
625/625 [=====] - 76s 121ms/step - loss: 0.0933 -
accuracy: 0.9712 - val_loss: 0.1253 - val_accuracy: 0.9770
Epoch 5/10
625/625 [=====] - 76s 121ms/step - loss: 0.0838 -
accuracy: 0.9723 - val_loss: 0.1100 - val_accuracy: 0.9750
Epoch 6/10
625/625 [=====] - 76s 121ms/step - loss: 0.0800 -
accuracy: 0.9756 - val_loss: 0.1241 - val_accuracy: 0.9770
Epoch 7/10
625/625 [=====] - 78s 124ms/step - loss: 0.0750 -
accuracy: 0.9764 - val_loss: 0.1200 - val_accuracy: 0.9750
Epoch 8/10
625/625 [=====] - 76s 121ms/step - loss: 0.0723 -
accuracy: 0.9771 - val_loss: 0.1492 - val_accuracy: 0.9730
Epoch 9/10
625/625 [=====] - 76s 121ms/step - loss: 0.0685 -
accuracy: 0.9790 - val_loss: 0.1232 - val_accuracy: 0.9690
Epoch 10/10
625/625 [=====] - 78s 124ms/step - loss: 0.0735 -
accuracy: 0.9786 - val_loss: 0.1499 - val_accuracy: 0.9750
```

```
[95]: import matplotlib.pyplot as plt
accuracy = history.history["accuracy"]
val_accuracy = history.history["val_accuracy"]
loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(accuracy) + 1)
plt.plot(epochs, accuracy, "bo", label="Training accuracy")
plt.plot(epochs, val_accuracy, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "ro", label="Training loss")
plt.plot(epochs, val_loss, "r", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
plt.show()
```





```
[97]: # Load the saved model
test_model_9 = keras.models.load_model("fine_tuning3.h5")

# Evaluate the model on the test dataset
test_loss, test_acc = test_model_9.evaluate(test_dataset_4)

# Print the test accuracy
print(f"Test accuracy: {test_acc:.3f}")
```

32/32 [=====] - 3s 94ms/step - loss: 0.1100 - accuracy: 0.9750

Test accuracy: 0.975

SCRATCH MODELS:

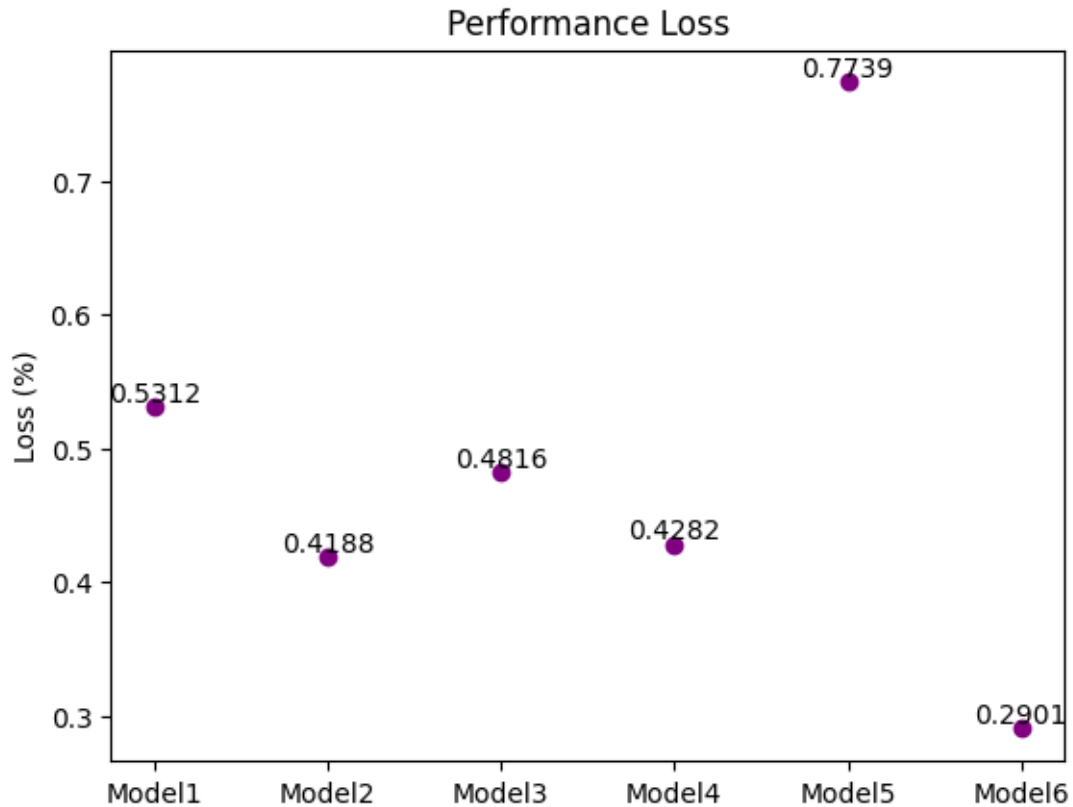
```
[98]: loss_dict = {'Model11': 0.5312, 'Model12': 0.4188, 'Model13': 0.4816, 'Model14': 0.4282,
                'Model15': 0.7739, 'Model16': 0.2901}

# Get model names and loss values as separate lists
models = list(loss_dict.keys())
losses = list(loss_dict.values())
```



```
# Plot the scatter plot with labels
plt.scatter(models, losses, color='purple')
plt.title('Performance Loss')
plt.ylabel('Loss (%)')

for (xi, yi) in zip(models, losses):
    plt.text(xi, yi, str(yi), va='bottom', ha='center')
```



```
[99]: import matplotlib.pyplot as plt

# Create a dictionary with model names as keys and accuracy values as values
acc_dict = {'Model1': 0.729, 'Model2': 0.801, 'Model3': 0.772, 'Model4': 0.801,
            'Model5': 0.871, 'Model6': 0.897}


# Get model names and accuracy values as separate lists
models = list(acc_dict.keys())
accuracy = list(acc_dict.values())

# Plot the scatter plot with labels
plt.scatter(models, accuracy, color='red')
```

```
plt.title('Performance Accuracy')
plt.ylabel('Accuracy (%)')

for (xi, yi) in zip(models, accuracy):
    plt.text(xi, yi, str(yi), va='bottom', ha='center')

plt.show()
```



PRE-TRAINED MODELS:

```
[100]: loss_dict = {'Model17': 0.1688, 'Model18': 0.0844, 'Model19': 0.0346}

# Get model names and loss values as separate lists
models = list(loss_dict.keys())
losses = list(loss_dict.values())


# Plot the scatter plot with labels
plt.scatter(models, losses, color='purple')
plt.title('Performance Loss')
plt.ylabel('Loss (%)')
```

```

for (xi, yi) in zip(models, losses):
    plt.text(xi, yi, str(yi), va='bottom', ha='center')

plt.show()

```



```


[101]: # Create a dictionary with model names as keys and accuracy values as values
acc_dict = {'Model7': 0.975, 'Model8': 0.981, 'Model9': 0.991}

# Get model names and accuracy values as separate lists
models = list(acc_dict.keys())
accuracy = list(acc_dict.values())

# Plot the scatter plot with labels
plt.scatter(models, accuracy, color='red')
plt.title('Performance Accuracy')
plt.ylabel('Accuracy (%)')

for (xi, yi) in zip(models, accuracy):
    plt.text(xi, yi, str(yi), va='bottom', ha='center')

```



[ ]: